

Forecasting

Forecasting without a model

Forecasting without a model: Plan

- Converting time series into cross-sectional data
- Add lags to the y_t variable
- Use aggregation and sliding or growing windows

Forecasting

Adding predictors

There are algorithms that, based on the training sample of the dependent variable y , learning matrix of predictors X , and new predictors X_F build a forecast \hat{y}_F

Random Forest, gradient boosting... and even linear regression!

You can average ARIMA/ETS forecasts and forecasts from other algorithms

How to create predictors?

From one column y you can create an entire matrix of predictors X !

- Use **lags** y_{t-k}
- Use **functions of lags** as predictors

Using y lags

For example, let's take two lags, Ly_t and L^2y_t

Training sample:

$$\begin{pmatrix} y_3 \\ y_4 \\ y_5 \\ \vdots \\ y_T \end{pmatrix} \quad \begin{pmatrix} y_1 & y_2 \\ y_2 & y_3 \\ y_3 & y_4 \\ \vdots & \vdots \\ y_{T-2} & y_{T-1} \end{pmatrix}$$

Sample for prediction:

$$\begin{pmatrix} ? \end{pmatrix} \quad \begin{pmatrix} y_{T-1} & y_T \end{pmatrix}$$

How many lags to add?

- Each added lag **reduces** the training sample!
- It is reasonable to add **closest lags** Ly_t, L^2y_t
- For seasonal data it is reasonable to add a **seasonal lag** $L^{12}y_t$
- Some algorithms are **sensitive to extra predictors** (e.g. regression)
- Some algorithms are **insensitive to extra predictors** (e.g. a random forest)

Functions of lags

When predicting y_t we can use any function of **previous** y_{t-1} , y_{t-2} , ...lags

For example:

- $\Delta y_{t-1} = y_{t-1} - y_{t-2}$;
- $\max\{y_{t-1}, y_{t-2}, y_{t-3}\}$;
- $\min\{y_{t-1}, y_{t-2}, \dots, y_1\}$.

Typical Predictor

- **Aggregate function:**
Min, Max, Mean, Median, Range, Sample Variance, Sample Standard Deviation, ...
- **Sliding Window:** The aggregate function can be applied to, say, the previous three values $y_{t-1}, y_{t-2}, y_{t-3}$.
- **Growing Window:** The aggregate function can be applied to all previous values $y_{t-1}, y_{t-2}, \dots, y_1$.

Using y lag functions

For example, let's take the maximum as a sliding window and the minimum as a growing window

Training sample:

$$\begin{pmatrix} y_3 \\ y_4 \\ y_5 \\ \vdots \\ y_T \end{pmatrix} \quad \begin{pmatrix} \max\{y_1, y_2\} & \min\{y_1, y_2\} \\ \max\{y_2, y_3\} & \min\{y_1, y_2, y_3\} \\ \max\{y_3, y_4\} & \min\{y_1, \dots, y_4\} \\ \vdots & \vdots \\ \max\{y_{T-2}, y_{T-1}\} & \min\{y_1, \dots, y_{T-1}\} \end{pmatrix}$$

Sample for prediction:

$$\begin{pmatrix} ? \end{pmatrix} \quad \begin{pmatrix} \max\{y_{T-1}, y_T\} & \min\{y_1, \dots, y_T\} \end{pmatrix}$$

Forecasting without a model: Summary

- Remember about random forest, gradient boosting and even regular regression
- Add dependent variable lags
- Add aggregation functions as a sliding and growing window

More predictors!

More predictors: Plan

- Trend predictors
- Seasonal and holiday dummy
- Cosines and sines as predictors

Let's use the time!

Let's take t and \sqrt{t} as an example

Training sample:

$$\begin{pmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_T \end{pmatrix} \quad \begin{pmatrix} 1 & \sqrt{1} \\ 2 & \sqrt{2} \\ 3 & \sqrt{3} \\ \vdots & \vdots \\ T & \sqrt{T} \end{pmatrix}$$

Sample for prediction:

$$\begin{pmatrix} ? \end{pmatrix} \quad \begin{pmatrix} T + 1 & \sqrt{T + 1} \end{pmatrix}$$

Monotonic transformations of time

- You can always **try!**
- For algorithms based on **decision trees** (random forest, gradient boosting) additional monotonic time transformations are **useless**
- Be aware of the possible transformation **of the original variable** (logarithm, Box-Cox transformation)

Seasonal and holiday dummy

If there are **not many** seasons, then it is reasonable to include a dummy for each season

Training sample for quarterly data:

$$\begin{pmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \\ y_5 \\ y_6 \\ \vdots \\ y_T \end{pmatrix} \quad \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

The dummy variable trap

In **regression**, be aware of the dummy variable **trap**!

- Either use a dummy for every season and a model without a constant,
- or use a dummy for all seasons except one and a model with a constant

Algorithms based on **decision trees** (random forest, gradient boosting) are **resistant** to the dummy variable trap

Why do we need sines and cosines?

Adding all dummy variables works **poorly** if you need **a lot** of them.

It is hardly worth adding 365 dummy variables for **daily** data.

Sine and cosine will help to decrease the **number** of predictors!

Two facts:

- The period of $\sin t$ and $\cos t$ is 2π
- Multiplying the argument by a reduces the **period** by a factor of a

Fourier expansion

Theorem

Any continuous and differentiable function f with period 2π can be represented as

$$f(t) = c + \sum_{k=1}^{\infty} a_k \cos(kt) + b_k \sin(kt)$$

Practical recipe for daily data:

- Add predictors $\cos\left(\frac{2\pi}{365} \cdot t\right)$ and $\sin\left(\frac{2\pi}{365} \cdot t\right)$
- Add predictors $\cos\left(\frac{2\pi}{365} \cdot 2t\right)$ and $\sin\left(\frac{2\pi}{365} \cdot 2t\right)$
- Add predictors $\cos\left(\frac{2\pi}{365} \cdot 3t\right)$ and $\sin\left(\frac{2\pi}{365} \cdot 3t\right)$
- ...

More predictors: Summary

- Use **time** as a predictor
- Seasonality in predictors can be reflected using **dummy variables** or using **cosine** and **sine** functions

Predictors and *ARIMA*

Predictors and *ARIMA*: Plan

- Regression with *ARMA* errors
- *ARMAX* model
- *ARDL* model

Regression with *ARMA* errors

Equation

$$y_t = \beta_1 + \beta_2 a_t + \beta_3 b_t + \varepsilon_t,$$

where a_t and b_t are **predictors**

- The series $(y_t), (a_t), (b_t), (\varepsilon_t)$ **stationary**
- $\varepsilon_t \sim ARMA(p, q)$ with respect to white noise (u_t)
- $\mathbb{E}(\varepsilon_t \mid a_{t-1}, b_{t-1}, a_{t-2}, b_{t-2}, \dots) = 0$
- **Fourth moments** of predictors are finite

ARMAX model

Equation

$$y_t = c + \gamma_1 a_t + \gamma_2 b_t + \beta_1 y_{t-1} + \dots + \beta_p y_{t-p} + u_t + \alpha_1 u_{t-1} + \dots + \alpha_q u_{t-q},$$

where a_t and b_t are **predictors** and (u_t) is **white noise**

- Series (y_t) , (a_t) , (b_t) **stationary**
- $\mathbb{E}(u_t \mid a_{t-1}, b_{t-1}, y_{t-1}, a_{t-2}, b_{t-2}, y_{t-2}, \dots) = 0$
- **The fourth predictor moments** are finite

ARMAX model is not completely equivalent to regression with ARMA errors, but gives **approximately the same** quality of predictions

Properties of the *ARMAX* model and regression with *ARMA* errors

- If the model assumptions aren't violated, then the maximum likelihood estimators **are consistent**
- For non-stationary variables (y_t) and predictors (a_t) and (b_t) , we can switch to the first **differences**
- Estimators **remain** consistent if trend, seasonality dummy and trigonometric predictors are added
- **Not any** predictor makes it possible to obtain a consistent estimator of the coefficient
- **Sometimes** you can get good predictions even if the assumptions are violated

ARDL model

ARDL — AutoRegressive Distributed Lag model

Autoregressive model with distributed lags

The $ARDL(p, q)$ model equation

$$y_t = c + \beta_1 y_{t-1} + \dots + \beta_p y_{t-p} + x_t + \alpha_1 x_{t-1} + \dots + \alpha_q x_{t-q} + u_t$$

- (u_t) errors are **white noise**
- Process (x_t) **or** process (Δx_t) is stationary
- Process (y_t) is **non-stationary**, but (Δy_t) is stationary
- $\mathbb{E}(u_t \mid y_{t-1}, x_{t-1}, y_{t-2}, x_{t-2}, \dots) = 0$

Properties of the *ARDL* model

- Using **predictor lags** (x_t) instead of noise lags (u_t)
- Suitable for **non-stationary** (y_t)
- Used to find a **long-term relationships** between time series
- If the model assumptions aren't violated, then **the OLS estimators are consistent**, although they are biased
- You can add **multiple predictors** with different number of lags

Predictors and *ARIMA*: Summary

- For **stationary data** you can use regression with *ARMA* errors or *ARMAX* model
- Regression with *ARMA* errors can be constructed for the **differences**
- For **non-stationary series** it is sometimes possible to use the *ARDL* model

Model Quality

Model Quality: Plan

- Scale-based metrics
- Percentage-based metrics

Remember the goal!

If the goal of building a model is forecasts one step ahead, then it is reasonable to compare models in predictive strength one step ahead.

If the goal is to detect the moment of model discord, then it is reasonable to look for a model that gives the minimum error when there is no discord, and the maximum error when there is discord.

Notations for brevity

For the forecast, it is important **when** it is built, and for **how many steps ahead**:

$$\hat{y}_{t+h|t}$$

Sometimes for **short**:

$$\hat{y}_{t+h}$$

Problem:

$$\hat{y}_{(t+1)+2} \neq \hat{y}_{(t+2)+1}$$

Anti-quality metrics

Forecast error: $e_{t+h} = y_{t+h} - \hat{y}_{t+h}$

Mean Absolute Error:

$$MAE = \frac{|e_{T+1}| + |e_{T+2}| + \dots + |e_{T+H}|}{H}$$

Root Mean Squared Error:

$$RMSE = \sqrt{\frac{e_{T+1}^2 + e_{T+2}^2 + \dots + e_{T+H}^2}{H}}$$

Scaling

Convert error $e_{t+h} = y_{t+h} - \hat{y}_{t+h}$ to percentage $p_t = e_t/y_t \cdot 100$
or $p_t^s = e_t/(0.5y_t + 0.5\hat{y}_t) \cdot 100$

Mean Absolute Percentage Error:

$$MAPE = \frac{|p_{T+1}| + |p_{T+2}| + \dots + |p_{T+H}|}{H}$$

Symmetric Mean Absolute Percentage Error:

$$sMAPE = \frac{|p_{T+1}^s| + |p_{T+2}^s| + \dots + |p_{T+H}^s|}{H}$$

Automatically compare with naive model

Naive: $\hat{y}_t^{naive} = y_{t-1}$ or $\hat{y}_t^{naive} = y_{t-12}$ Let's scale our forecast error e_t to MAE^{naive} :

$$q_t = \frac{e_t}{MAE^{naive}}$$

Mean Absolute Scaled Error:

$$MASE = \frac{|q_{T+1}| + |q_{T+2}| + \dots + |q_{T+H}|}{H}$$

Comparing q to 1 compares our model with the naive one

Model Quality: Summary

- MAE, RMSE
- MAPE
- MASE

Model Comparison

Model Comparison: Plan

- Cross-validation
- Akaike criterion

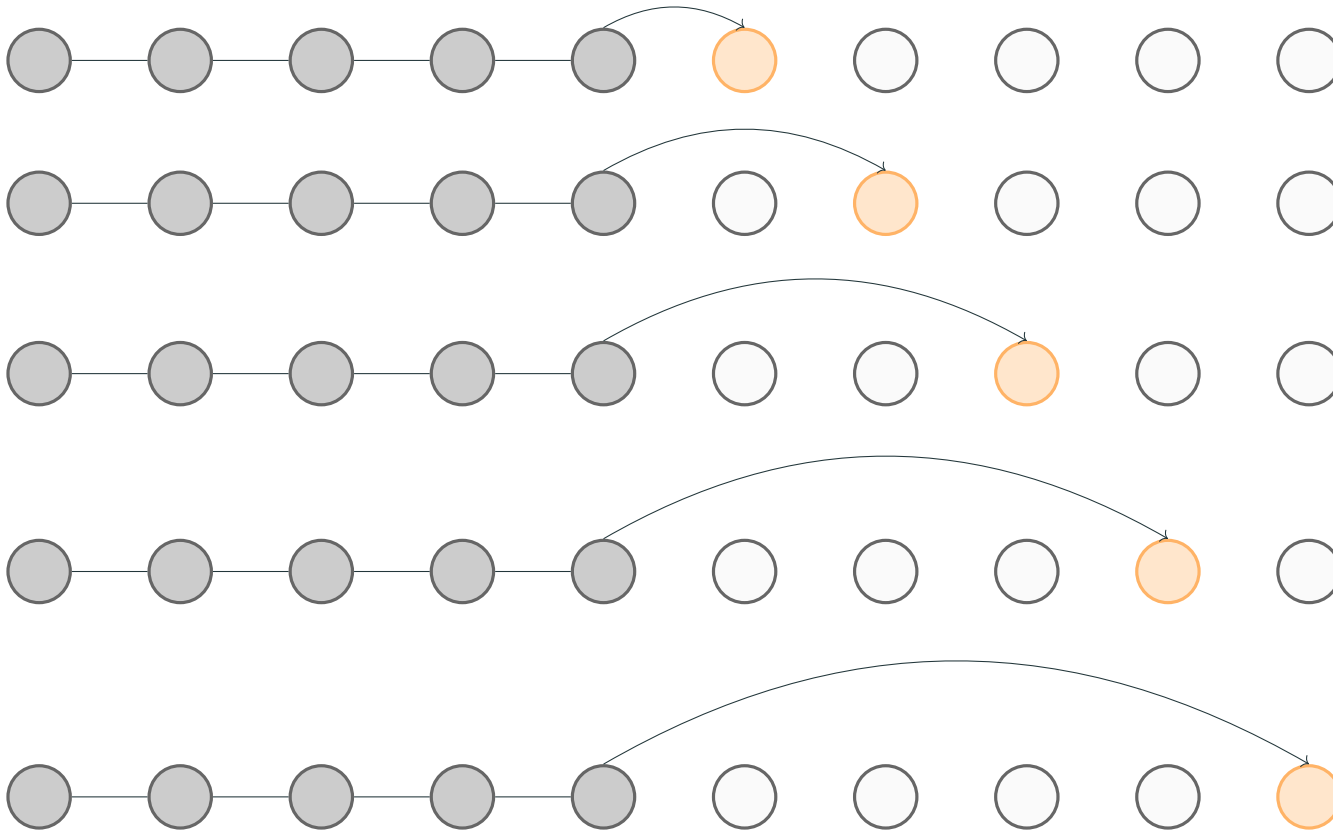
Train and test set

Strategy:

1. Split the entire sample into **training** (at the beginning) and **test** (at the end) sets
2. **Evaluate** several models on the training set
3. **Predict** each observation of the test sample using each model
4. Calculate prediction **errors**
5. **Compare** models by MAE and choose the best one

Disadvantage: **forecasts have different horizons**

Dividing into train and test




Sliding window Cross-validation

Strategy:

1. Select the starting size for **train** sample (at the beginning)
2. **Evaluate** several models on the train set
3. **Predict** one step ahead with each model
4. Calculate prediction **errors**
5. **Shift** the training sample one observation to the right
6. Repeat steps 2-5
7. **Compare** models by MAE and choose the best one

Sliding window Cross-validation




Growing window Cross-validation

Strategy:

1. Select the starting size for **train** sample (at the beginning)
2. Evaluate several models on the training set
3. Predict one step ahead with each model
4. Calculate prediction errors
5. **Increase** the training set by one observation.
6. Repeat steps 2-5
7. Compare models by MAE and choose the best one

Growing window Cross-validation



Cross-validation Discussion

Sliding window cross-validation: there are many observations and we suspect that dependencies between values can change.

Growing window cross-validation: there are few observations or we are sure that the dependency persists.

Cross-validation can be **time consuming**!

Let's make cross-validation quicker!

Approximate cross-validation by one step forward based on $RMSE$ using... **Akaike Information Criterion:**

$$AIC = -2 \ln L + 2k,$$

where $\ln L$ is the logarithm of the maximum likelihood on the training set, k is the total number of model parameters

Nuances of AIC

- AIC has theoretical grounds:

$$\frac{AIC_A - AIC_B}{2} \approx KL(\text{Truth} || \text{Model A}) - KL(\text{Truth} || \text{Model B})$$

- May be used for non-nested models
- For Gaussian y_t models, the criterion approximates comparison over $RMSE$
- The models being compared must be for the same observations

Model Comparison: Summary

- Cross-validation: sliding and growing window
- AIC is a fast and approximate analogue of cross-validation

Forecast comparison

Forecast comparison: Plan

- Diebold-Mariano test
- Test **assumptions**
- Test implementation

Diebold-Mariano Test

- Used to compare **two** forecasts
- Compares the forecasts for the specified **forecast horizon** h
- Not optimal for **model comparison**
- Not suitable for **pairwise** comparison of multiple forecasts

DM test assumptions

Consider difference of two **forecast losses**:

$$d_t = e_{A,t}^2 - e_{B,t}^2, \quad e_{\text{Model},t} = \hat{y}_{\text{Model},t} - y_t$$

Difference d_t assumed to be **stationary**:

$$\mathbb{E}(d_t) = \mu_d,$$

$$\text{Cov}(d_t, d_{t-k}) = \gamma_k,$$

in particular,

$$\text{Var}(d_t) = \gamma_0$$

Test method

Under the correct $H_0 : \mu_d = 0$:

$$DM = \frac{\bar{d}}{se(\bar{d})} \rightarrow \mathcal{N}(0; 1),$$

where $se^2(\bar{d})$ is a consistent estimator for $\text{Var}(\bar{d})$

In practice, we evaluate the regression on a constant

$$\hat{d}_t = \hat{\beta}_1,$$

get $\hat{\beta}_1 = \bar{d}$ and use **robust standard errors**,

$$DM = \frac{\hat{\beta}_1}{se_{HAC}(\hat{\beta}_1)}.$$

How does robust estimators work?

Compare forecasts for P steps ahead,

$$\text{Var}(\bar{d}) = \frac{\text{Var}(d_1) + \text{Var}(d_2) + \dots + 2 \text{Cov}(d_1, d_2) + \dots}{P^2}$$

From the stationarity of d_t , the variance of $\text{Var}(\bar{d})$ is

$$\frac{P\gamma_0 + 2(P-1)\gamma_1 + 2(P-2)\gamma_2 + \dots}{P^2}$$

The naive estimator for the variance of $\widehat{\text{Var}}(\bar{d})$ is

$$\frac{P\hat{\gamma}_0 + 2(P-1)\hat{\gamma}_1 + 2(P-2)\hat{\gamma}_2 + \dots}{P^2}$$

Why compare forecasts?

Nuance: comparison of forecasts and comparison of models are different tasks.

A model can win a lot **in simplicity** and lose a little in predictions.

On a small sample **loss of information** about the quality of forecasts on the training sample is significant.

Forecast comparison: Summary

- The Diebold-Mariano test is suitable for comparing **two** forecasts
- Comparing forecasts and comparing models are **slightly different** tasks